

Engineering Cloud Computing

Assignment 1

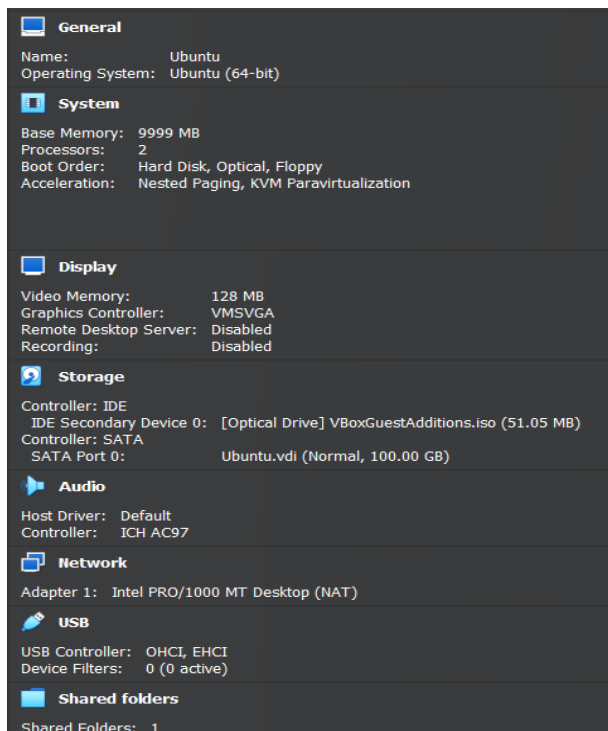
Dilip Nikhil Francies

dfrancie@iu.edu

2001274638

Note : Due to limited availability of Jetstream2 instance, I have installed and configured Hadoop on my Windows laptop, following the tutorial uploaded on Canvas.

Here are the specifications of my virtual machine:



Part 1:

- Output the top-3 IP addresses with the granularity of an hour
 - a) You should read the two examples.
 - b) Develop your code based on examples. The program may take more than one round of MapReduce.

Part 1 – Mapper Function:

```
#Dilip Nikhil Francies
#Part 1 mapper function

import sys
import re
from datetime import datetime

# Regular expression pattern to match IP addresses and timestamps
log_pattern = re.compile(r'(?P<ip>\d+\.\d+\.\d+\.\d+) .*'
\[(?P<timestamp>[^\]]+)\]')

for line in sys.stdin:
    match = log_pattern.search(line.strip()) #find matches in the input
    if match:
        ip_address = match.group('ip')
        timestamp_str = match.group('timestamp').split()[0] # Extracting
date from timestamp
        try:
            timestamp = datetime.strptime(timestamp_str,
"%d/%b/%Y:%H:%M:%S") #strip the time into Hour, minutes and seconds
            hour = timestamp.strftime("%H")
            print(f"{hour}\t{ip_address}\t1") # print the hour and the
associated IP address
        except ValueError:
            pass
```

Part 1 – Reducer Function:

```
#Author: Dilip Nikhil Francies
#reducer function

import sys
from operator import itemgetter
from collections import defaultdict
from datetime import datetime, timedelta

hourly_ip_count = defaultdict(lambda: defaultdict(int))
for line in sys.stdin: #input
    line = line.strip()
    hour, ip, count = line.split('\t', 2)
    count = int(count)
    hourly_ip_count[hour][ip] += count

# Iterate through each hour
for hour, ip_counts in hourly_ip_count.items():
    start_time = datetime.strptime(hour, "%H")
    end_time = start_time + timedelta(hours=1) - timedelta(seconds=1) #
from 00:00:00 to 00:00:59 every hour
    hour_range = f"{start_time.strftime('%H:%M:%S')} to
{end_time.strftime('%H:%M:%S')}"
    # Sort IP addresses based on their counts
    sorted_ips = sorted(ip_counts.items(), key=itemgetter(1), reverse=True)
    # Output the top 3 IP addresses for each hour
    top_3_ips = sorted_ips[:3] #once sorted, get the top three IPs
    for ip, count in top_3_ips:
        print(f"From {hour_range}, IP: {ip}, Count: {count}") #print
```

Steps followed:

1. Copy log files and the source code from the localhost to the VM through a shared folder, and copy it to the “Hadoop” user from “Dilip”.

```
dilip@Ubuntu:~/Desktop$ sudo cp -r ECC /home/hadoop/
dilip@Ubuntu:~/Desktop$ su - hadoop
Password:
hadoop@Ubuntu:~$ ls
ECC  hadoop  hadoop-3.3.6.tar.gz  hadoopdata  snap
hadoop@Ubuntu:~$ cd ECC
hadoop@Ubuntu:~/ECC$ ls
a1      logstat1_reduce.py  logstat2_reduce.py
logstat1_map.py  logstat2_map.py    sample.log
```

2. Run “Start-all.sh” to initiate the startup process of various Hadoop daemons.

```
hadoop@Ubuntu:~/ECC$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [Ubuntu]
Starting resourcemanager
Starting nodemanagers
```

3. Convert access.log to hdfs [Hadoop distributed file system] and start the map reduce job.

MapReduce code:

```
mapred streaming
-input /access.log
-file logstat1_map.py
-mapper "python3 logstat1_map.py"
-file logstat1_reduce.py
-reducer "python3 logstat1_reduce.py"
-output /sample_output_access
```

```
hadoop@Ubuntu:~/ECC$ hdfs dfs -copyFromLocal access.log /
hadoop@Ubuntu:~/ECC$ mapred streaming -input /access.log -file logstat1_map.py -mapper "python3 logstat1_map.py" -file logstat1_reduce.py -reducer "python3
logstat1_reduce.py" -output /sample_output_access
2024-03-05 17:04:47,972 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [logstat1_map.py, logstat1_reduce.py] [/home/hadoop/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.3.6.jar] /tmp/streamjob7045267247551393
445.jar tmpDir=null
2024-03-05 17:04:51,844 INFO client.DefaultNoHARMAFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2024-03-05 17:04:52,813 INFO client.DefaultNoHARMAFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2024-03-05 17:04:53,660 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1709660894826_0
007
2024-03-05 17:04:54,657 INFO mapred.FileInputFormat: Total input files to process : 1
2024-03-05 17:04:54,715 INFO net.NetworkTopology: Adding a new node: /default-rack/127.0.0.1:9866
2024-03-05 17:04:54,852 INFO mapreduce.JobSubmitter: number of splits:26
2024-03-05 17:04:55,407 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1709660894826_0007
2024-03-05 17:04:55,408 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-03-05 17:04:55,963 INFO conf.Configuration: resource-types.xml not found
2024-03-05 17:04:55,964 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-03-05 17:04:56,181 INFO impl.YarnClientImpl: Submitted application application_1709660894826_0007
2024-03-05 17:04:56,355 INFO mapreduce.Job: The url to track the job: http://Ubuntu.myguest.virtualbox.org:8088/proxy/application_1709660894826_0007/
2024-03-05 17:04:56,358 INFO mapreduce.Job: Running job: job_1709660894826_0007
2024-03-05 17:05:09,898 INFO mapreduce.Job: Job job_1709660894826_0007 running in uber mode : false
2024-03-05 17:05:09,900 INFO mapreduce.Job: map 0% reduce 0%
2024-03-05 17:05:57,013 INFO mapreduce.Job: map 1% reduce 0%
```

4. Once the MapReduce job is initiated, make sure it terminates with a successful output generation and check for the files written in the folder /sample_output_access.

```
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=3502543223
File Output Format Counters
Bytes Written=4268
2024-03-05 17:13:05,056 INFO streaming.StreamJob: Output directory: /sample_output_access
hadoop@Ubuntu:~/ECC$ hadoop fs -ls /sample_output_access
Found 2 items
-rw-r--r-- 1 hadoop supergroup 0 2024-03-05 17:13 /sample_output_access/_SUCCESS
-rw-r--r-- 1 hadoop supergroup 4268 2024-03-05 17:13 /sample_output_access/part-00000
```

5. View the output:

Format: "From *Start time* to *End time*, IP: *XX.XXX.XX.XXX*, Count: *XXXX*"

```
hadoop@Ubuntu:~/ECC$ hdfs dfs -cat /sample_output_access/*
From 00:00:00 to 00:59:59, IP: 66.249.66.194, Count: 14298
From 00:00:00 to 00:59:59, IP: 66.249.66.91, Count: 12232
From 00:00:00 to 00:59:59, IP: 66.249.66.92, Count: 4291
From 01:00:00 to 01:59:59, IP: 66.249.66.91, Count: 13874
From 01:00:00 to 01:59:59, IP: 66.249.66.194, Count: 12485
From 01:00:00 to 01:59:59, IP: 66.249.66.92, Count: 2924
From 02:00:00 to 02:59:59, IP: 66.249.66.91, Count: 11697
From 02:00:00 to 02:59:59, IP: 66.249.66.194, Count: 10345
From 02:00:00 to 02:59:59, IP: 91.99.72.15, Count: 1448
From 03:00:00 to 03:59:59, IP: 66.249.66.194, Count: 8644
From 03:00:00 to 03:59:59, IP: 66.249.66.91, Count: 7914
From 03:00:00 to 03:59:59, IP: 91.99.72.15, Count: 1275
From 04:00:00 to 04:59:59, IP: 66.249.66.194, Count: 10805
From 04:00:00 to 04:59:59, IP: 66.249.66.91, Count: 7571
From 04:00:00 to 04:59:59, IP: 91.99.72.15, Count: 1511
From 05:00:00 to 05:59:59, IP: 66.249.66.194, Count: 10534
From 05:00:00 to 05:59:59, IP: 66.249.66.91, Count: 7035
From 05:00:00 to 05:59:59, IP: 91.99.72.15, Count: 1921
```

Figure: Top three IPs for every Hour.

Note: Full-length output is provided in the logs folder with the file named "part1_output.txt"

Part 2:

1. Make your program like a database search.
 - a. Your program should be able to accept parameters from users, such as 0-1, which means from time 00:00 to 01:00, and output the top 3 IP addresses in the given time period.

For Part 2 of the assignment, I developed a shell program that lets the user enter the time frame that he would like to query in the format "*start_time*[space]*end_time*". These values are then passed on to my MapReduce code as arguments to run the job.

```
#!/usr/bin/env bash
# Prompt the user to enter two numbers separated by a space
echo "Enter start and end time separated by a space:"
read start end

# Check if both inputs are valid integers
if ! [[ "$start" =~ ^[0-9]+$ ]] || ! [[ "$end" =~ ^[0-9]+$ ]]; then
    echo "Error: Input is not valid. Please enter two integers separated by a space."
    exit 1
fi

mapred streaming -input /access.log -output /part2_A -file
logstat2_map.py -mapper "python3 logstat2_map.py $start $end" -file
logstat2_reduce.py -reducer "python3 logstat2_reduce.py"
```

Part 2 Mapper Function:

```
#Author : Dilip Nikhil Francies
#mapper function

import re
import sys
# Compile the regex pattern outside the loop for better performance
ip_hour_pattern =
re.compile(r'(?P<ip>\d+\.\d+\.\d+\.\d+).*?\d{4}:(?P<hour>\d{2}):\d{2}')
# Retrieve command-line arguments directly

try:
    start_period, end_period = map(int, sys.argv[1:3]) # pars from
arguments
except (ValueError, IndexError):
    print("Invalid input. Please provide a valid range as command-line
arguments (e.g., '3 4')")
    sys.exit(1)

for line in sys.stdin: #read from acces.log
    line = line.strip()
    match = ip_hour_pattern.search(line)
    if match:
        ip_address = match.group('ip')
        hour = int(match.group('hour'))
        if start_period <= hour < end_period: # check if start time is less
than end time
            print(f"[{hour:02}:00] {ip_address} 1") #print ip and hour.
```

Part 2 Reducer function:

```
#author : Dilip Nikhil Francies
#reducer function for part 2

import sys
# Create an empty dictionary to store IP counts
ip_counts = {}

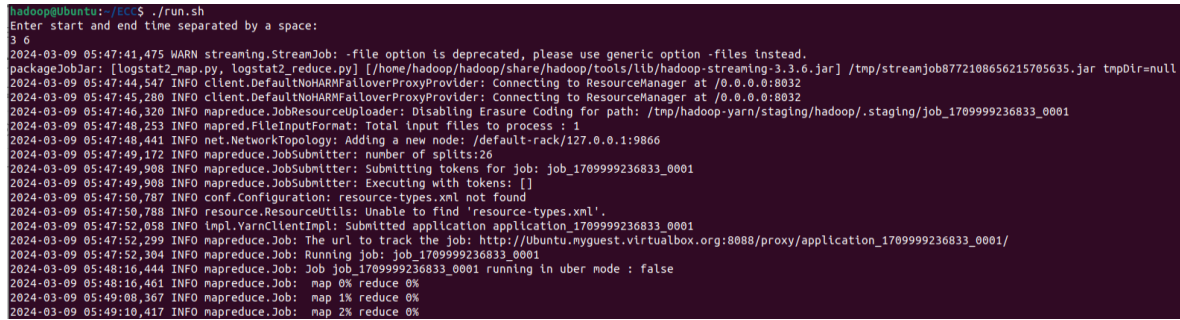
for line in sys.stdin:
    line = line.strip()
    if line.startswith("["):
        # Extract hour and IP address
        hour, ip = line.split()[0], line.split()[1]
        # Strip brackets from hour
        hour = hour.strip("[")
        # Update IP count in the dictionary
        ip_counts[ip] = ip_counts.get(ip, 0) + 1

sorted_ip_counts = sorted(ip_counts.items(), key=lambda x: x[1],
reverse=True)
# Print the dictionary containing IP counts
for ip, count in sorted_ip_counts[:3]: #sort top three
    print(f"[{hour}] {ip} {count}")
```

Running the shell file “./run.sh” first prompts the user to input a time frame, in this example it would be from 3AM to 6 AM. These inputs will be passed on to my mapper function with arguments as \$start \$end.

MapReduce code used in the shell:

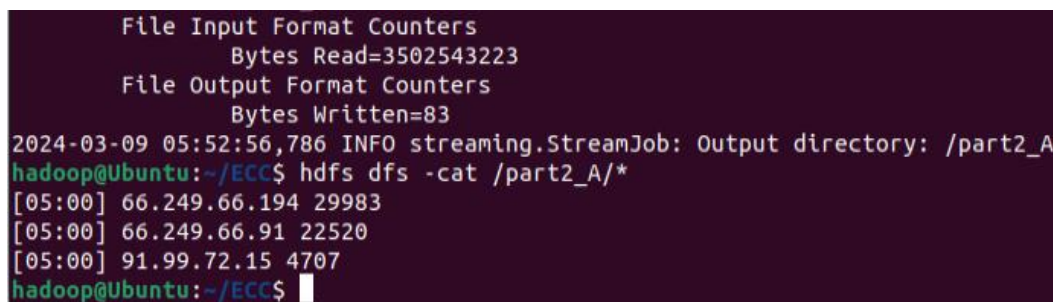
```
mapred streaming
-input /access.log
-output /part2_A
-file logstat2_map.py
-mapper "python3 logstat2_map.py $start $end"
-file logstat2_reduce.py
-reducer "python3 logstat2_reduce.py"
```



```
hadoop@Ubuntu:~/ECC$ ./run.sh
Enter start and end time separated by a space:
3 6
2024-03-09 05:47:41,475 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [logstat2_map.py, logstat2_reduce.py] [/home/hadoop/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.3.6.jar] /tmp/streamjob8772108656215705635.jar tmpDir=null
2024-03-09 05:47:44,547 INFO client.DefaultHadoopHARMFalloverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2024-03-09 05:47:45,280 INFO client.DefaultHadoopHARMFalloverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2024-03-09 05:47:46,320 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1709999236833_0001
2024-03-09 05:47:48,253 INFO mapred.FileInputFormat: Total input files to process : 1
2024-03-09 05:47:48,441 INFO net.NetworkTopology: Adding a new node: /default-rack/127.0.0.1:9866
2024-03-09 05:47:49,172 INFO mapreduce.JobSubmitter: number of splits:26
2024-03-09 05:47:49,908 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1709999236833_0001
2024-03-09 05:47:49,908 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-03-09 05:47:50,787 INFO conf.Configuration: resource-types.xml not found
2024-03-09 05:47:50,788 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-03-09 05:47:52,050 INFO impl.VarClientImpl: Submitted application application_1709999236833_0001
2024-03-09 05:47:52,299 INFO mapreduce.Job: The url to track the job: http://Ubuntu.nyquest.virtualbox.org:8088/proxy/application_1709999236833_0001/
2024-03-09 05:47:52,304 INFO mapreduce.Job: Running job: job_1709999236833_0001
2024-03-09 05:48:16,444 INFO mapreduce.Job: Job job_1709999236833_0001 running in uber mode : false
2024-03-09 05:48:16,461 INFO mapreduce.Job: map 0% reduce 0%
2024-03-09 05:49:08,367 INFO mapreduce.Job: map 1% reduce 0%
2024-03-09 05:49:10,417 INFO mapreduce.Job: map 2% reduce 0%
```

Figure: Initiating the MapReduce program

Once the job is run, the output is saved in /part2_A folder which then can be viewed as shown below:



```
File Input Format Counters
  Bytes Read=3502543223
File Output Format Counters
  Bytes Written=83
2024-03-09 05:52:56,786 INFO streaming.StreamJob: Output directory: /part2_A
hadoop@Ubuntu:~/ECC$ hdfs dfs -cat /part2_A/*
[05:00] 66.249.66.194 29983
[05:00] 66.249.66.91 22520
[05:00] 91.99.72.15 4707
hadoop@Ubuntu:~/ECC$
```

Figure: Part 2 Output of top three IPs in the given time frame

Part 2:

- b. Run it along with three other examples, WordCount, Sort, Grep, at the same time, and test fair and capacity schedulers.

A Brief Introduction to Schedulers:

Scheduling mainly refers to the process of allocating and managing computing resources among multiple jobs and users in a distributed computing environment to ensure efficient resource utilization and fair sharing of resources among different applications running on a Hadoop cluster.

Scheduling is essential in Hadoop for proper resource management, efficient utilization by keeping all nodes busy while minimizing resource wastage, prioritizing different jobs based on importance, urgency, and finally failure tolerance to maintain jobs in progress and to reduce disruptions

Types of Schedulers:

Fair Scheduler - The Fair Scheduler in Hadoop is a resource scheduler designed to ensure fair and equitable allocation of computing resources among users and jobs in a Hadoop cluster. Overall, the Fair Scheduler promotes fairness, optimizes resource utilization, and ensures predictable performance for all users and jobs in a Hadoop cluster.

Capacity Scheduler - The Capacity Scheduler in Hadoop is a resource scheduler that facilitates efficient and predictable resource allocation. Key features include flexible and efficient mechanisms for managing resources in shared Hadoop clusters, ensuring fair access, predictable performance, and effective utilization.

As for part 2, question B, the following steps that I performed to test different schedulers and their effect on resource allocation to Hadoop jobs.

1. Configure the "yarn-site.xml" to indicate the preferred scheduler to be used.

```
<configuration>
  <property>
    <name>yarn.resourcemanager.scheduler.class</name>
    <value>org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.CapacityScheduler</value>
  </property>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
</configuration>
```

Fig: Capacity scheduler configuration

```
<configuration>
  <property>
    <name>yarn.resourcemanager.scheduler.class</name>
    <value>org.apache.hadoop.yarn.server.resourcemanager.scheduler.fair.FairScheduler</value>
  </property>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
</configuration>
```

Fig: Fair scheduler configuration

2. Once the configuration is complete, stop and restart the resource manager daemon with the following command

```
hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon stop resourcemanager
hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon start resourcemanager
```

3. Open 4 different terminals and execute 4 MapReduce programs parallelly

1. Top 3 IPs,
2. WordCount,
3. Grep
4. Sort

```
hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon stop resourcemanager
hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon start resourcemanager

hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon stop resourcemanager
hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon start resourcemanager

hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon stop resourcemanager
hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon start resourcemanager

hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon stop resourcemanager
hadoop@Ubuntu:~/hadoop/etc/hadoop$ yarn --daemon start resourcemanager
```

These are the following commands that I used:

1. Top 3 IP:

```
mapred streaming
-input /access.log
-output /part2_B_B
-file logstat2_map.py
-mapper "python3 logstat2_map.py $start $end"
-file logstat2_reduce.py
-reducer "python3 logstat2_reduce.py"
```

2. WordCount:

```
mapred streaming
-input /access.log
-output /wc2
-mapper /bin/cat
-reducer /bin/wc
```

3. Sort

```
mapred streaming -input /sample.log
-output /sort
-mapper /bin/sort
-reducer /bin/cat
```

4. Grep

```
mapred streaming
-input /access.log
-output /grep2
-mapper "/bin/grep '66.249.66.194'"
-reducer /bin/cat
```

Discussion:

Capacity Scheduler:

When the 4 jobs are run, all the resources are allocated to the first process in the queue and all the users can use the maximum resources available in the cluster. Initially, the first program is allocated with 7 CPU vcores and 8192 of memory, which is the entire resource available in the cluster.

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Allocated GPUs
application 1710027705392_0002	hadoop	streamjob4556474596164970945.jar	MAPREDUCE		default	0	Sat Mar 9 18:46:02 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1
application 1710027705392_0002	hadoop	streamjob5807326587261156015.jar	MAPREDUCE		default	0	Sat Mar 9 18:45:55 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1
application 1710027705392_0001	hadoop	streamjob5352048026640018783.jar	MAPREDUCE		default	0	Sat Mar 9 18:45:33 -0500 2024	Sat Mar 9 18:45:38 -0500 2024	N/A	RUNNING	UNDEFINED	7	7	8192	-1

Once the 1st program is successfully executed complete, the resource utilization drops.

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Allocated GPUs	Reserved CPU V-Cores	Reserved Memory MB	Reserved GPUs	% of Queue	% of Cluster	Progress	Time
application 1710027705392_0003	hadoop	streamjob3184170254857409358.jar	MAPREDUCE		default	0	Sat Mar 9 18:46:45 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1	0	0	-1	0.0	0.0		80s
application 1710027705392_0003	hadoop	streamjob4556474596164970945.jar	MAPREDUCE		default	0	Sat Mar 9 18:46:02 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1	0	0	-1	0.0	0.0		80s
application 1710027705392_0002	hadoop	streamjob5807326587261156015.jar	MAPREDUCE		default	0	Sat Mar 9 18:45:55 -0500 2024	Sat Mar 9 18:51:55 -0500 2024	N/A	ACCEPTED	UNDEFINED	1	1	2048	-1	0	0	-1	25.0	25.0		80s

At this point the second process in queued and the resource allocations shoots up again, to all the available resources.

ID	User	Name	Application Type	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Allocated GPUs	Reserved CPU V-Cores	Reserved Memory MB	Reserved GPUs	% of Queue	% of Cluster	Progress
application-1710077703362-0004	hadoop	streamjob10841702548574097058.jar	MAPREDUCE		default	0	Sat Mar 9 18:48:43 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1	0	0	-1	0.0	0.0	
application-1710077703362-0003	hadoop	streamjob435647439616497045.jar	MAPREDUCE		default	0	Sat Mar 9 18:48:02 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1	0	0	-1	0.0	0.0	
application-1710077703362-0002	hadoop	streamjob5807320587261156015.jar	MAPREDUCE		default	0	Sat Mar 9 18:45:55 -0500 2024	Sat Mar 9 18:51:55 -0500 2024	N/A	RUNNING	UNDEFINED	7	7	8192	-1	0	0	-1	100.0	100.0	

The same process continues as long as all the jobs are complete in sequential manner, providing all the available resources to jobs that are executed in priority.

Application Queues

Legend

Capacity

Used

Used lower capacity

Max Capacity

Users Respecting Resources

Auto Created Queues

1

Application Queues

37.5% used

37.5% used

Show 25 entries

ID	User	Name	Application Type	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Allocated GPUs	Reserved CPU V-Cores	Reserved Memory MB	Reserved GPUs	% of Queue	% of Cluster	Progress
application-1710077703362-0004	hadoop	streamjob10841702548574097058.jar	MAPREDUCE		default	0	Sat Mar 9 18:48:43 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1	0	0	-1	0.0	0.0	
application-1710077703362-0003	hadoop	streamjob435647439616497045.jar	MAPREDUCE		default	0	Sat Mar 9 18:48:02 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1	0	0	-1	0.0	0.0	
application-1710077703362-0002	hadoop	streamjob5807320587261156015.jar	MAPREDUCE		default	0	Sat Mar 9 18:45:55 -0500 2024	Sat Mar 9 18:51:55 -0500 2024	N/A	RUNNING	UNDEFINED	2	2	3072	-1	0	0	-1	37.5	37.5	

Show 25 entries

ID	User	Name	Application Type	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Allocated GPUs	Reserved CPU V-Cores	Reserved Memory MB	Reserved GPUs	% of Queue	% of Cluster	Progress
application-1710077703362-0004	hadoop	streamjob10841702548574097058.jar	MAPREDUCE		default	0	Sat Mar 9 18:48:43 -0500 2024	Sat Mar 9 19:06:54 -0500 2024	Sat Mar 9 19:07:38 -0500 2024	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	
application-1710077703362-0003	hadoop	streamjob435647439616497045.jar	MAPREDUCE		default	0	Sat Mar 9 18:48:02 -0500 2024	Sat Mar 9 19:02:34 -0500 2024	Sat Mar 9 19:06:46 -0500 2024	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	
application-1710077703362-0002	hadoop	streamjob5807320587261156015.jar	MAPREDUCE		default	0	Sat Mar 9 18:45:55 -0500 2024	Sat Mar 9 18:51:55 -0500 2024	Sat Mar 9 19:02:28 -0500 2024	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	
application-1710077703362-0001	hadoop	streamjob535504802640018763.jar	MAPREDUCE		default	0	Sat Mar 9 18:45:38 -0500 2024	Sat Mar 9 18:45:38 -0500 2024	Sat Mar 9 18:51:47 -0500 2024	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	

Note: The output of each program is available in the folder attached.

Fair scheduler:

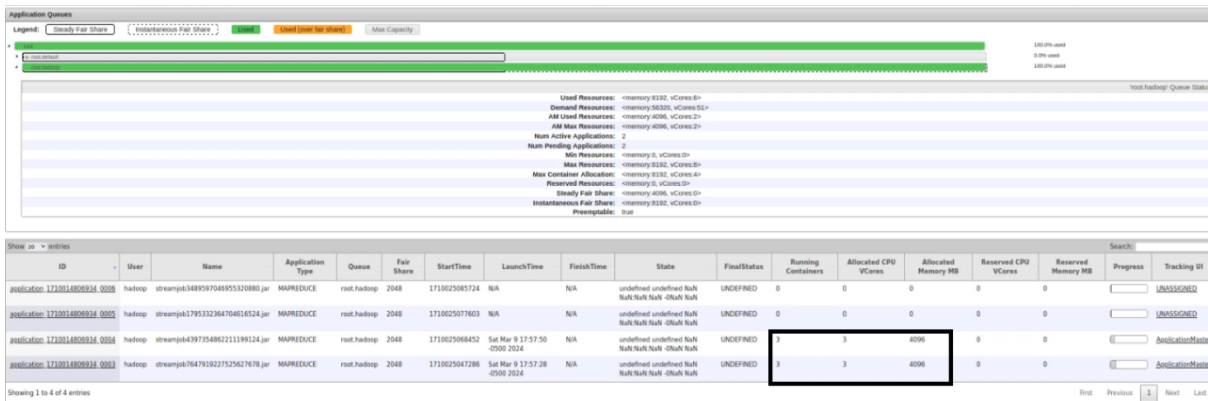
Fair scheduler on the other hand distributes the resources equally to the jobs that demand resources. Users get equal amount of resources [Dilip 50% Hadoop 50% [default value] as displayed below

ID	User	Name	Application Type	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Allocated GPUs	Reserved CPU V-Cores	Reserved Memory MB	Reserved GPUs	% of Queue	% of Cluster	Progress
application-1710077703362-0004	hadoop	streamjob10841702548574097058.jar	MAPREDUCE		default	0	Sat Mar 9 18:48:43 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1	0	0	-1	0.0	0.0	
application-1710077703362-0003	hadoop	streamjob435647439616497045.jar	MAPREDUCE		default	0	Sat Mar 9 18:48:02 -0500 2024	N/A	N/A	ACCEPTED	UNDEFINED	0	0	0	-1	0	0	-1	0.0	0.0	
application-1710077703362-0002	hadoop	streamjob5807320587261156015.jar	MAPREDUCE		default	0	Sat Mar 9 18:45:55 -0500 2024	Sat Mar 9 18:51:55 -0500 2024	N/A	RUNNING	UNDEFINED	2	2	3072	-1	0	0	-1	37.5	37.5	


Based on available resource, it is distributed to two jobs received.

ID	User	Name	Application Type	Queue	Fair Share	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Reserved CPU V-Cores	Reserved Memory MB	Progress	Tracking UI
application-1710077703362-0004	hadoop	streamjob10841702548574097058.jar	MAPREDUCE	root.hadoop	2048	1710025085724	N/A	N/A	undefined undefined NaN NaN NaN NaN	UNDEFINED	0	0	0	0	0		UNASSIGNED
application-1710077703362-0003	hadoop	streamjob1799332364704016124.jar	MAPREDUCE	root.hadoop	2048	1710025077603	N/A	N/A	undefined undefined NaN NaN NaN NaN	UNDEFINED	0	0	0	0	0		UNASSIGNED
application-1710077703362-0002	hadoop	streamjob4397354862211199124.jar	MAPREDUCE	root.hadoop	2048	1710025088452	Sat Mar 9 17:57:50 -0500 2024	N/A	undefined undefined NaN NaN NaN NaN	UNDEFINED	1	1	2048	0	0		UNASSIGNED
application-1710077703362-0001	hadoop	streamjob7847931027525627678.jar	MAPREDUCE	root.hadoop	2048	1710025047096	Sat Mar 9 17:57:28 -0500 2024	N/A	undefined undefined NaN NaN NaN NaN	UNDEFINED	1	1	2048	0	0		UNASSIGNED

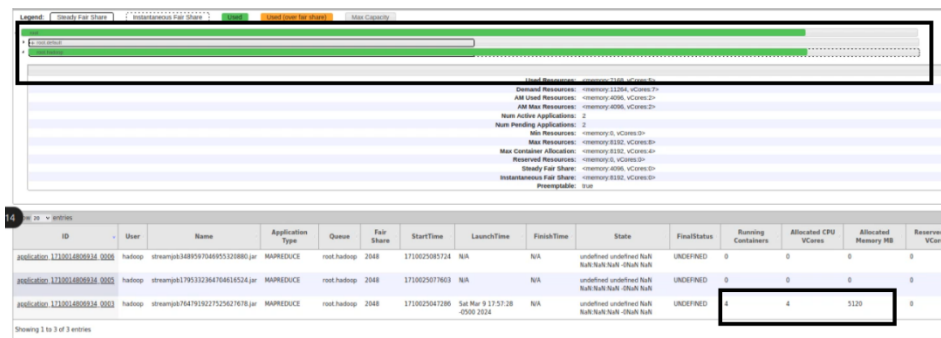
Because, the other user “Dilip” has no jobs running, and the resources are available, the “Hadoop” user makes use of it and distributes available resources to jobs 1 and 2.



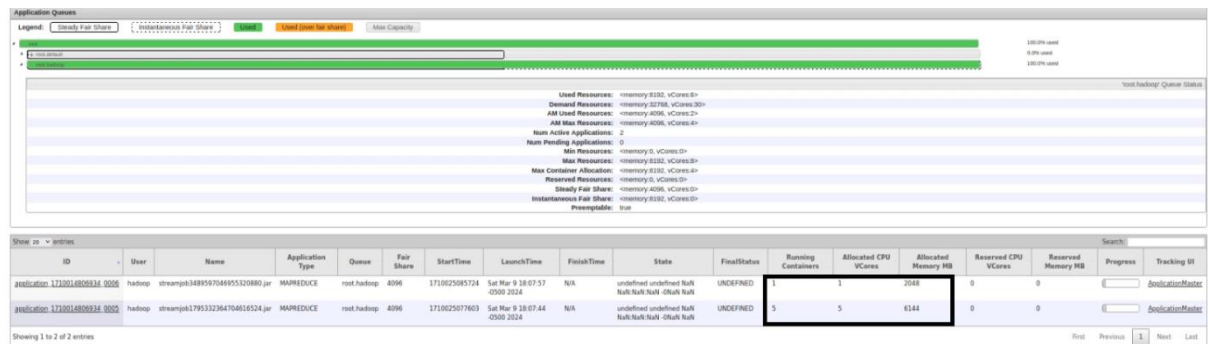
As suggested by TA Juliette, I used “top” command to see resources utilization of the VM.



Once the first job in complete the remaining resources are allocated to the jobs that were running earlier, ensuring that all the jobs have some number of resources allocated. At the same time, once the previous job is completed, the remaining is distributed equally among other processes.



After the second job is complete, the remaining resources are allocated to 3rd and 4th job.



Now the jobs are distributed equally:

PID	USER	PR	NI	VIRT	RES	SHR	S	PCPU	PMEM	TIME+	COMMAND	
25032	hadoop	20	0	2483968	301136	28800	S	46.1	3.0	0:18.32	java	ob: map 8% reduce 0%
25398	hadoop	20	0	2448908	98684	28416	S	36.5	1.0	0:01.84	java	ob: map 9% reduce 0%
25288	hadoop	20	0	2475040	193704	28800	S	29.9	2.0	0:07.93	java	ob: map 10% reduce 0%
23212	hadoop	20	0	2465492	179620	28800	S	28.6	1.8	0:07.97	java	ob: map 11% reduce 0%
25374	hadoop	20	0	2445528	82216	28288	S	22.4	0.8	0:01.59	java	ob: map 14% reduce 0%
25285	hadoop	20	0	2464908	129444	28672	S	21.1	1.3	0:05.79	java	ob: map 15% reduce 0%
17442	hadoop	20	0	4323208	987824	29952	S	6.2	9.9	2:05.22	java	ob: map 19% reduce 0%
3016	dillip	20	0	4747008	558840	187436	S	2.0	5.6	32:30.65	gnome-shell	ob: map 20% reduce 0%
24838	hadoop	20	0	2812852	418988	29056	S	2.0	4.2	0:24.33	java	ob: map 21% reduce 0%
17331	hadoop	20	0	4482644	646128	29952	S	1.0	5.5	1:59.36	java	ob: map 22% reduce 0%
20999	dillip	20	0	21876	4224	3328	R	0.7	0.0	0:02.50	top	ob: map 23% reduce 0%
16	root	20	0	0	0	0	R	0.3	0.0	0:20.12	rcu_preempt	ob: map 25% reduce 0%
23	root	20	0	0	0	0	R	0.3	0.0	0:06.32	ksoftirqd/1	ob: map 26% reduce 0%
606	root	20	0	1318984	30632	19968	S	0.3	0.3	0:00.56	snmpd	
3553	dillip	20	0	227804	3456	2944	S	0.3	0.0	0:22.12	VBoxClient	
3624	dillip	20	0	586832	67152	48804	S	0.3	0.7	1:08.92	gnome-terminal	
16697	hadoop	20	0	4265228	423044	29184	S	0.3	4.3	0:45.38	java	

Max Resources: <memory 8192, vCores 2>

Max Container Allocation: <memory 8192, vCores 4>

Reserved Resources: <memory 0, vCores 0>

Steady Fair Share: <memory 4096, vCores 0>

Instantaneous Fair Share: <memory 8192, vCores 0>

Preemptible: true

ID	User	Name	Application Type	Queue	Fair Share	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Reserved CPU V-Cores	Reserved Memory MB	Progress	Tracking UI
application: 17101480934 000	hadoop	streamjob1489197045532080.jar	MAPREDUCE	root.hadoop	4096	1710025087724	Sat Mar 9 18:07:57 CEST 2024	N/A	undefined undefined Null Null Null Null Null Null	UNDEFINED	3	3	4096	0	0	<div></div>	ApplicationMaster
application: 17101480934 000	hadoop	streamjob1795132364704616524.jar	MAPREDUCE	root.hadoop	4096	1710025077803	Sat Mar 9 18:07:44 CEST 2024	N/A	undefined undefined Null Null Null Null Null Null	UNDEFINED	3	3	4096	0	0	<div></div>	ApplicationMaster

Showing 1 to 2 of 2 entries

First Previous 1 Next Last

Final job process gets all the available resources:

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority	Scheduler Bug %
Fair Scheduler	[memory-mb (unit=MB), vcores]	<memory 1024, vCores 1>	<memory 8192, vCores 4>	0	0

Application Queues

Legend: Steady Fair Share, Instantaneous Fair Share, Used, Used (over fair share), Max Capacity

Used Resources: <memory 8192, vCores 7>

Demand Resources: <memory 2208, vCores 21>

AM Used Resources: <memory 2048, vCores 1>

AM Max Resources: <memory 4096, vCores 4>

Now Active Applications: 1

Min Resources: <memory 0, vCores 0>

Max Resources: <memory 8192, vCores 0>

Max Container Allocation: <memory 8192, vCores 4>

Reserved Resources: <memory 0, vCores 0>

Steady Fair Share: <memory 4096, vCores 0>

Instantaneous Fair Share: <memory 8192, vCores 0>

Preemptible: true

ID	User	Name	Application Type	Queue	Fair Share	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Reserved CPU V-Cores	Reserved Memory MB	Progress	Tracking UI
application: 17101480934 000	hadoop	streamjob1795132364704616524.jar	MAPREDUCE	root.hadoop	8192	1710025077803	Sat Mar 9 18:07:44 CEST 2024	N/A	undefined undefined Null Null Null Null Null Null	UNDEFINED	7	7	8192	0	0	<div></div>	ApplicationMaster

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

One thing that I have observed is, the number of resources allocated reduces when most of the mapping part is complete, which kind of makes sense assuming the mapper is usually process heavy in some of these programs, and only the required amount of resources allocated to that particular job

Comparison of different performance metrics for different types of schedulers:

Capacity										
	Program	total time -map tasks	total time - reduce tasks	MB ms map tasks	MB ms - reduce tasks	GC time elapsed	CPU time spent	Time Elapsed	Agg resource alloc	Agg vcore (s)
	ip	1708570	181746	1749575680	186107904	20604	211820	6:14	2769186	2311
	word_count	2302033	515957	2357281792	528339968	22598	545240	16:30	4251706	3499
	grep	1166133	134193	1194120192	137413632	17657	153520	20:43	1912387	1593
	sort	31123	6758	31869952	6920192	635	4660	6:14	2769186	2311
Fair										
	Program	total time -map tasks	total time - reduce tasks	MB ms map tasks	MB ms - reduce tasks	GC time elapsed	CPU time spent	Time Elapsed	Agg resource alloc	Agg vcore (s)
	ip	636525	239254	651801600	244996096	8565	97410	6:14	2769186	2311
	word_count	834627	520935	854658048	533437440	9501	392300	13:29	3091284	2187
	grep	799809	185287	819004416	189733888	13820	135190	14:26	1812821	1384
	sort	20848	7912	21348352	8101888	423	3890	14:51	159532	92

Fig: Various metrics compared between Fair and Capacity Schedulers

Graphs Plots:



Fig: Plots displaying performance metrics of Fair and Capacity Schedulers

Conclusion:

On the surface, fair scheduler definitely looks better in terms of overall performance, but I believe other parameters define what type of scheduler algorithms to use. Perhaps I should dive deeper to understand better and derive conclusions.

